

ISOT FLO
SUMMIT 2017



Fault tolerance

```
with tf.device(tf.train.replica_device_setter(ps_tasks=3)):  
    weights_1 = tf.get_variable("weights_1", [784, 100])  
    biases_1 = tf.get_variable("biases_1", [100])  
    # ...
```

```
saver = tf.train.Saver(sharded=True)
```

```
with tf.Session(server.target) as sess:  
    while True:  
        # ...
```

```
        if is_chief and step % 1000 == 0:  
            saver.save(sess, "gs://mrry/model/...")
```

Each PS task writes in parallel

Saver can write checkpoints
to a distributed file system



...

One worker task acts as "chief"



TensorFlow
DEV SUMMIT 2017

you can write to
Google Cloud Storage,

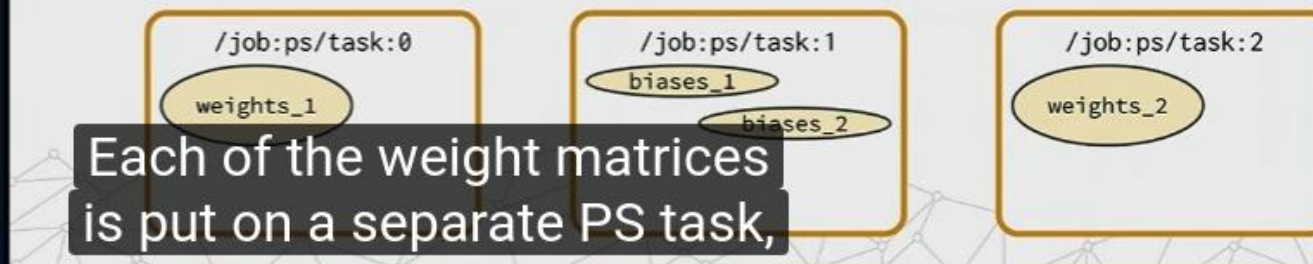
#tfdevsummit

Full screen



Load balancing variables

```
greedy = tf.contrib.training.GreedyLoadBalancingStrategy(...)
with tf.device(tf.train.replica_device_setter(
    ps_tasks=3, ps_strategy=greedy)):
    weights_1 = tf.get_variable("weights_1", [784, 100])
    biases_1 = tf.get_variable("biases_1", [100])
    weights_2 = tf.get_variable("weights_2", [100, 10])
    biases_2 = tf.get_variable("biases_2", [10])
```



Between-graph replication

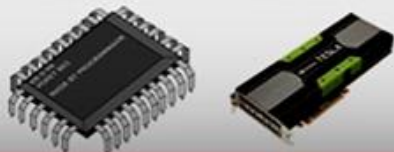
```
with tf.device("/job:ps/task:0/cpu:0"):
    W = tf.Variable(...)
    b = tf.Variable(...)
with tf.device("/job:worker/task:0/gpu:0"):
    output = tf.matmul(input, W) + b
    loss = f(output)
```

Client

/job:worker/task:0/

cpu:0

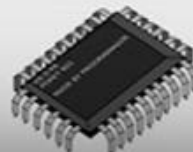
gpu:0



doing the same thing, with one
key difference in the device

/job:ps/task:0/

cpu:0



```
with tf.device("/job:ps/task:0/cpu:0"):
    W = tf.Variable(...)
    b = tf.Variable(...)
with tf.device("/job:worker/task:1/gpu:0"):
    output = tf.matmul(input, W) + b
    loss = f(output)
```

Client

/job:worker/task:1/

gpu:0

cpu:0



Full screen



10:26 / 28:06





Fault tolerance

MonitoredTrainingSession automates the recovery process

```
# Distributed code.  
server = tf.train.Server(...)  
is_chief = FLAGS.task_index == 0  
with tf.train.MonitoredTrainingSession(server.target, is_chief) as sess:  
    while not sess.should_stop():  
        sess.run(train_op)
```

Automatically initializes and/or
restores variables before returning



TensorFlow
DEV SUMMIT 2017

checkpoint if one is available,
before it returns control back

#tfdevsummit



Fault tolerance

```
with tf.device(tf.train.replica_device_setter(ps_tasks=3)):  
    weights_1 = tf.get_variable("weights_1", [784, 100])  
    biases_1 = tf.get_variable("biases_1", [100])  
    # ...
```

Each PS task writes in parallel

```
saver = tf.train.Saver(sharded=True)
```

```
with tf.Session(server.target) as sess:  
    while True:  
        # ...  
        if is_chief and step % 1000 == 0:  
            saver.save(sess, "/home/mrry/...")
```

One worker task acts as "chief"



at the start of day and logging
summaries for TensorBoard.

Sessions and Servers

Distributed TensorFlow runs on a *cluster* of servers

```
# Distributed code for a worker task
```

```
cluster = tf.train.ClusterSpec({"worker": ["192.168.0.1:2222", ...],  
                               "ps": ["192.168.1.1:2222", ...]})
```

cluster defines the set of processes

```
server = tf.train.Server(cluster, job_name="worker", task_index=0)
```

```
with tf.Session(server.target) as sess:
```

```
# ...
```

server represents a particular task in cluster

sess can run code on any device in cluster

Google Developers

Subscribed

27:21



Experiments and Estimators

High-level APIs package up the whole distributed workflow

```
def experiment_fn(config, params):  
    features = [tf.layers.embedding_column(...),  
                tf.layers.bucketized_column(...)]  
    return Experiment(  
        train_input_fn=..., eval_input_fn=...,  
        estimator=DNNClassifier(  
            hidden_units=[10, 20], feature_columns=features,  
            config, params))  
  
learn_runner.run(experiment_fn, config, ...)
```



uses a recently added class
called Experiment to package up



TensorFlow
DEV SUMMIT 2017

#tfdevsummit

Partitioned variables

```
greedy = tf.contrib.training.GreedyLoadBalancingStrategy(...)  
with tf.device(tf.train.replica_device_setter(  
    ps_tasks=3, ps_strategy=greedy)):
```

```
    embedding = tf.get_variable(  
        "embedding", [1000000000, 20],  
        partitioner=tf.fixed_size_partitioner(3))
```

here, TensorFlow will split
the large logical variable

/job:ps/task:0

embedding[0]

/job:ps/task:1

embedding[1]

/job:ps/task:2

embedding[2]



14:04 / 28:06

ISORTFIO
SUMM 2017



Sessions and Servers

Distributed TensorFlow runs on a *cluster* of servers

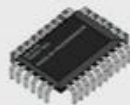
`tf.Session`

```
# Single-process code.  
with tf.Session() as sess:  
    sess.run(init_op)  
    for _ in range(NUM_STEPS):  
        sess.run(train_op)
```

`/job:worker/task:0/`

`cpu:0`

`gpu:0`



TensorFlow

that session will only
know about the devices

Google Developers

✓ Subscribed 1M

27:29

Sessions and Servers

Distributed TensorFlow runs on a *cluster* of servers

```
# Distributed code for a worker task.
```

```
cluster = tf.train.ClusterSpec({"worker": ["192.168.0.1:2222", ...],  
                               "ps": ["192.168.1.1:2222", ...]})
```

```
server = tf.train.Server(cluster, job_name="worker", task_index=0)
```

```
with tf.Session(server.target) as sess:
```

```
# ...
```

cluster defines the set of processes

server represents a
particular task in cluster

in that cluster.

27:49



Sessions and Servers

Distributed TensorFlow runs on a *cluster* of servers

```
# Distributed code for a PS task.  
cluster = tf.train.ClusterSpec({"worker": ["192.168.0.1:2222", ...],  
                               "ps": ["192.168.1.1:2222", ...]})  
  
server = tf.train.Server(cluster, job_name="ps", task_index=0)  
  
# Wait for incoming connections forever.  
server.join()
```



TensorFlow

PS task is much simpler.

In-graph replication

```
with tf.device("/job:ps/task:0/cpu:0"):
    W = tf.Variable(...)
    b = tf.Variable(...)
inputs = tf.split(0, num_workers, input)
outputs = []
for i in range(num_workers):
    with tf.device("/job:worker/task:%d/gpu:0" % i):
        outputs.append(tf.matmul(input[i], W) + b)
loss = f(outputs)
```

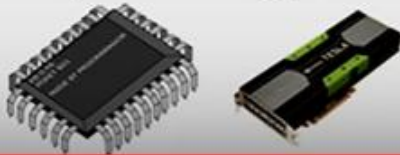
Client

like the earlier example.

/job:worker/task:0/

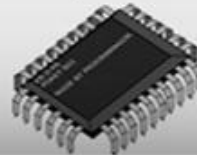
cpu:0

gpu:0



/job:ps/task:0/

cpu:0



/job:worker/task:1/

gpu:0

cpu:0



Default view

8:38 / 28:06

CC

HD

□

⌵

Experiments and Estimators

High-level APIs package up the whole distributed workflow

```
def experiment_fn(config, params):  
    features = [tf.layers.embedding_column(...),  
                tf.layers.bucketized_column(...)]  
    return Experiment(  
        train_input_fn=..., eval_input_fn=...,  
        estimator=DNNClassifier(  
            hidden_units=[10, 20], feature_columns=features,  
            config, params))
```

Declarative specification of a fully-connected neural network

```
learn_runner.run(experiment_fn, config)
```

You tell it how to read
a particular data set





Round-robin variables

```
with tf.device(tf.train.replica_device_setter(ps_tasks=3)):
```

```
weights_1 = tf.get_variable("weights_1", [784, 100])
```

```
biases_1 = tf.get_variable("biases_1", [100])
```

```
weights_2 = tf.get_variable("weights_2", [100, 10])
```

```
biases_2 = tf.get_variable("biases_2", [10])
```

/job:ps/task:0

weights_1

/job:ps/task:1

biases_1

/job:ps/task:2

weights_2

I guess I should have
drawn a diagram--



TensorFlow

Round-robin variables

```
with tf.device(tf.train.replica_device_setter(ps_tasks=3)):
```

```
weights_1 = tf.get_variable("weights_1", [784, 100])
```

```
biases_1 = tf.get_variable("biases_1", [100])
```

```
weights_2 = tf.get_variable("weights_2", [100, 10])
```

```
biases_2 = tf.get_variable("biases_2", [10])
```

placement strategies.

/job:ps/task:0

/job:ps/task:1



Sessions and Servers

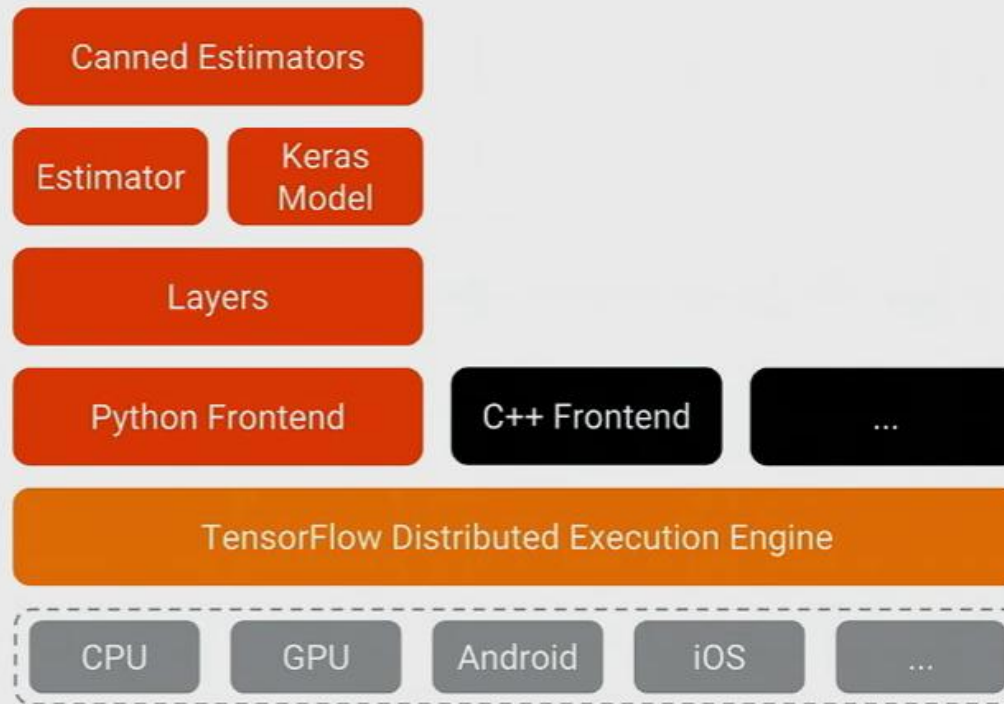
Distributed TensorFlow runs on a *cluster* of servers

```
# Distributed code for a worker task.  
cluster = tf.train.ClusterSpec({"worker": ["192.168.0.1:2222", ...],  
                               "ps": ["192.168.1.1:2222", ...]})  
  
server = tf.train.Server(cluster, job_name="worker", task_index=0)  
  
with tf.Session(server.target) as sess:  
    # ...
```

the first thing we need to do
is to provide a cluster spec.

26:04





TensorFlow

#tfdevsummit

Fault tolerance

```
with tf.device(tf.train.replica_device_setter(ps_tasks=3)):  
    weights_1 = tf.get_variable("weights_1", [784, 100])  
    biases_1 = tf.get_variable("biases_1", [100])  
    # ...
```

Each PS task writes in parallel

```
saver = tf.train.Saver(sharded=True)
```

```
with tf.Session(server.target) as sess:  
    while True:
```

```
        # ...
```

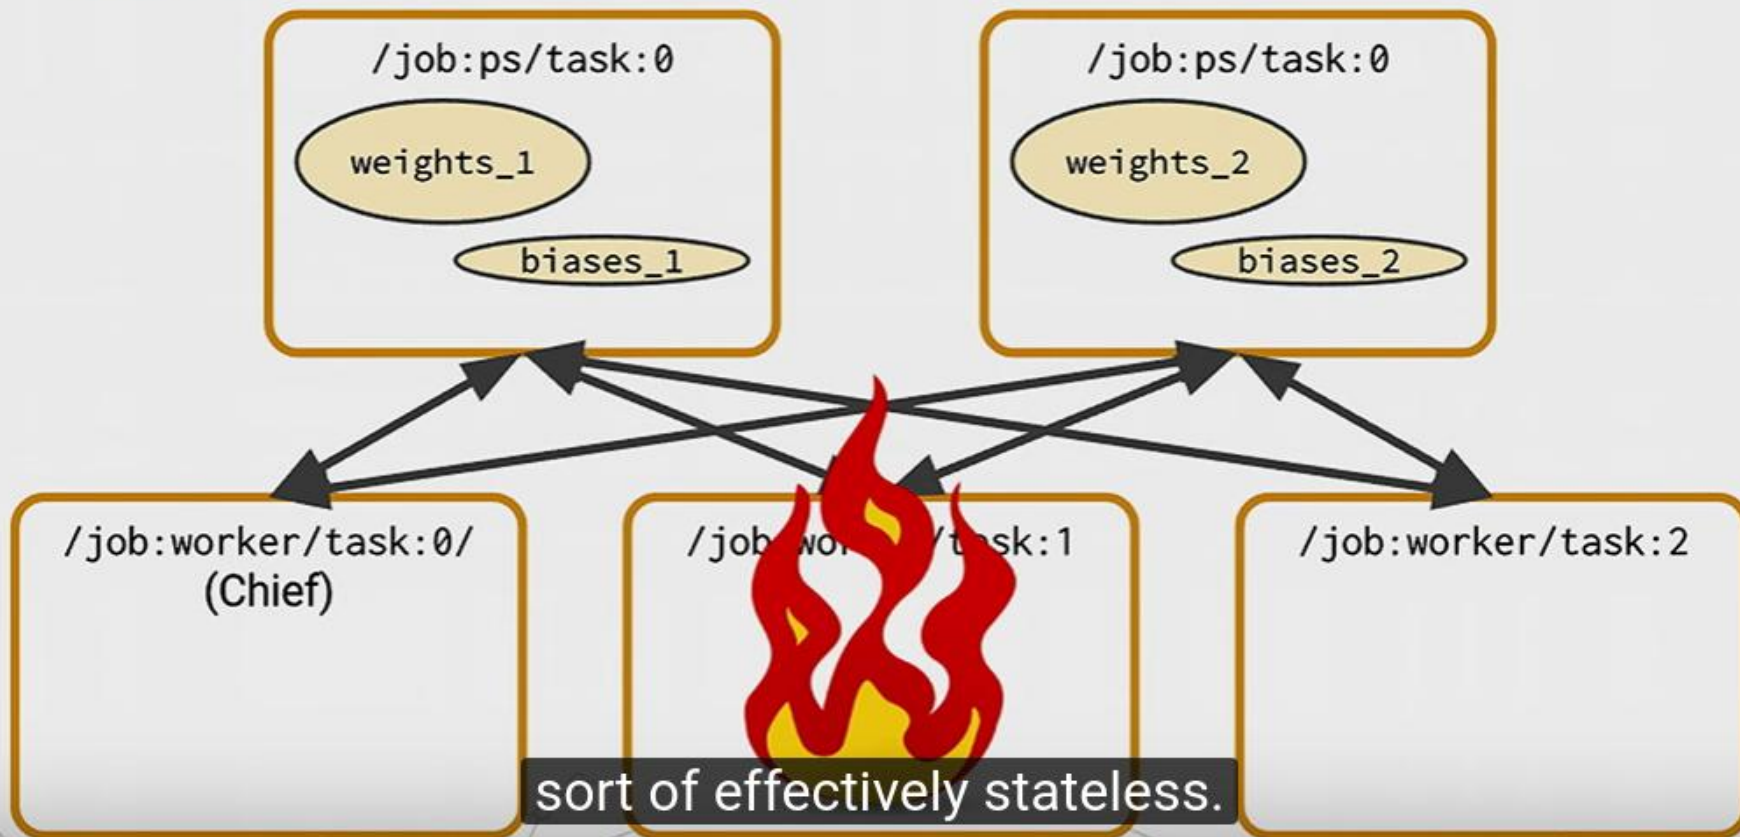
```
        if step % 1000 == 0:
```

```
            saver.save(sess, "model.ckpt")
```

want to set sharded equals true
when you create your saver.

26:40

Fault tolerance



Full screen



21:39 / 28:06



Fault tolerance

```
with tf.device(tf.train.replica_device_setter(ps_tasks=3)):  
    weights_1 = tf.get_variable("weights_1", [784, 100])  
    biases_1 = tf.get_variable("biases_1", [100])  
    # ...
```

```
saver = tf.train.Saver()
```

```
with tf.Session(server.target) as sess:  
    while True:
```

```
        # ...
```

```
        if step % 1000 == 0:
```

```
            saver.save(sess, "/home/mrry/...")
```

parameters to disk.

Google Developers

Subscribed

27:31

Fault tolerance

MonitoredTrainingSession automates the recovery process

```
# Distributed code.  
server = tf.train.Server(...)  
is_chief = FLAGS.task_index == 0  
with tf.train.MonitoredTrainingSession(server.target, is_chief) as sess:  
    while not sess.should_stop():  
        sess.run(train_op)
```

to know if it's
the chief or not,

Default view

23:57 / 28:06

Sessions and Servers

Distributed TensorFlow runs on a *cluster* of servers

Distributed code for a worker task.

```
cluster = tf.train.ClusterSpec({"worker": ["192.168.0.1:2222", ...],  
                               "ps": ["192.168.1.1:2222", ...]})
```

cluster defines the set of processes

```
server = tf.train.Server(cluster, job_name="worker", task_index=0)
```

```
with tf.Session(server.target) as sess:
```

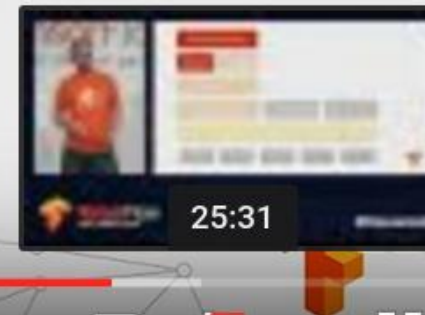
```
# ...
```

Fault tolerance

`MonitoredTrainingSession` automates the recovery process

```
# Single-process code.  
with tf.Session() as sess:  
    sess.run(init_op) # Or saver.restore(sess, ...)  
    for _ in range(NUM_STEPS):  
        sess.run(train_op)
```

before you start training.



Device placement

`tf.train.replica_device_setter()`

Simple heuristic for between-graph partitioning

- Round-robin variable placement by default
- Optional strategy for load balancing, partitioning
- All other ops placed on a worker task
- Customize using nested `with tf.device(...):` blocks

Variable placement

```
with tf.device("/job:ps/task:0"):  
  
    weights_1 = tf.get_variable("weights_1", [784, 100])  
    biases_1 = tf.get_variable("biases_1", [100])  
    weights_2 = tf.get_variable("weights_2", [100, 10])  
    biases_2 = tf.get_variable("biases_2", [10])
```

So far, I've just been
putting them on jobPS task 0

“

A distributed system is a system where I can't get my work done because a computer has failed that I've never even heard of.

”

— Leslie Lamport

on a set of machines, I hope you
take the wise words of Leslie

Google Developers

✓ Subscribed 1M

27:05